RSE Support for Georgia Tech's AI Makerspace

Dr. Fang (Cherry) Liu, Ronald Rahaman, Dr. Jeffrey Young

Senior Research Scientist/Manager of Research Software Engineer
PACE Research Computing Center at Georgia Tech

Presented at Research Software Engineers in HPC workshop at Supercomputing 2024
Atlanta, GA
Nov 17, 2024

Georgia Tech

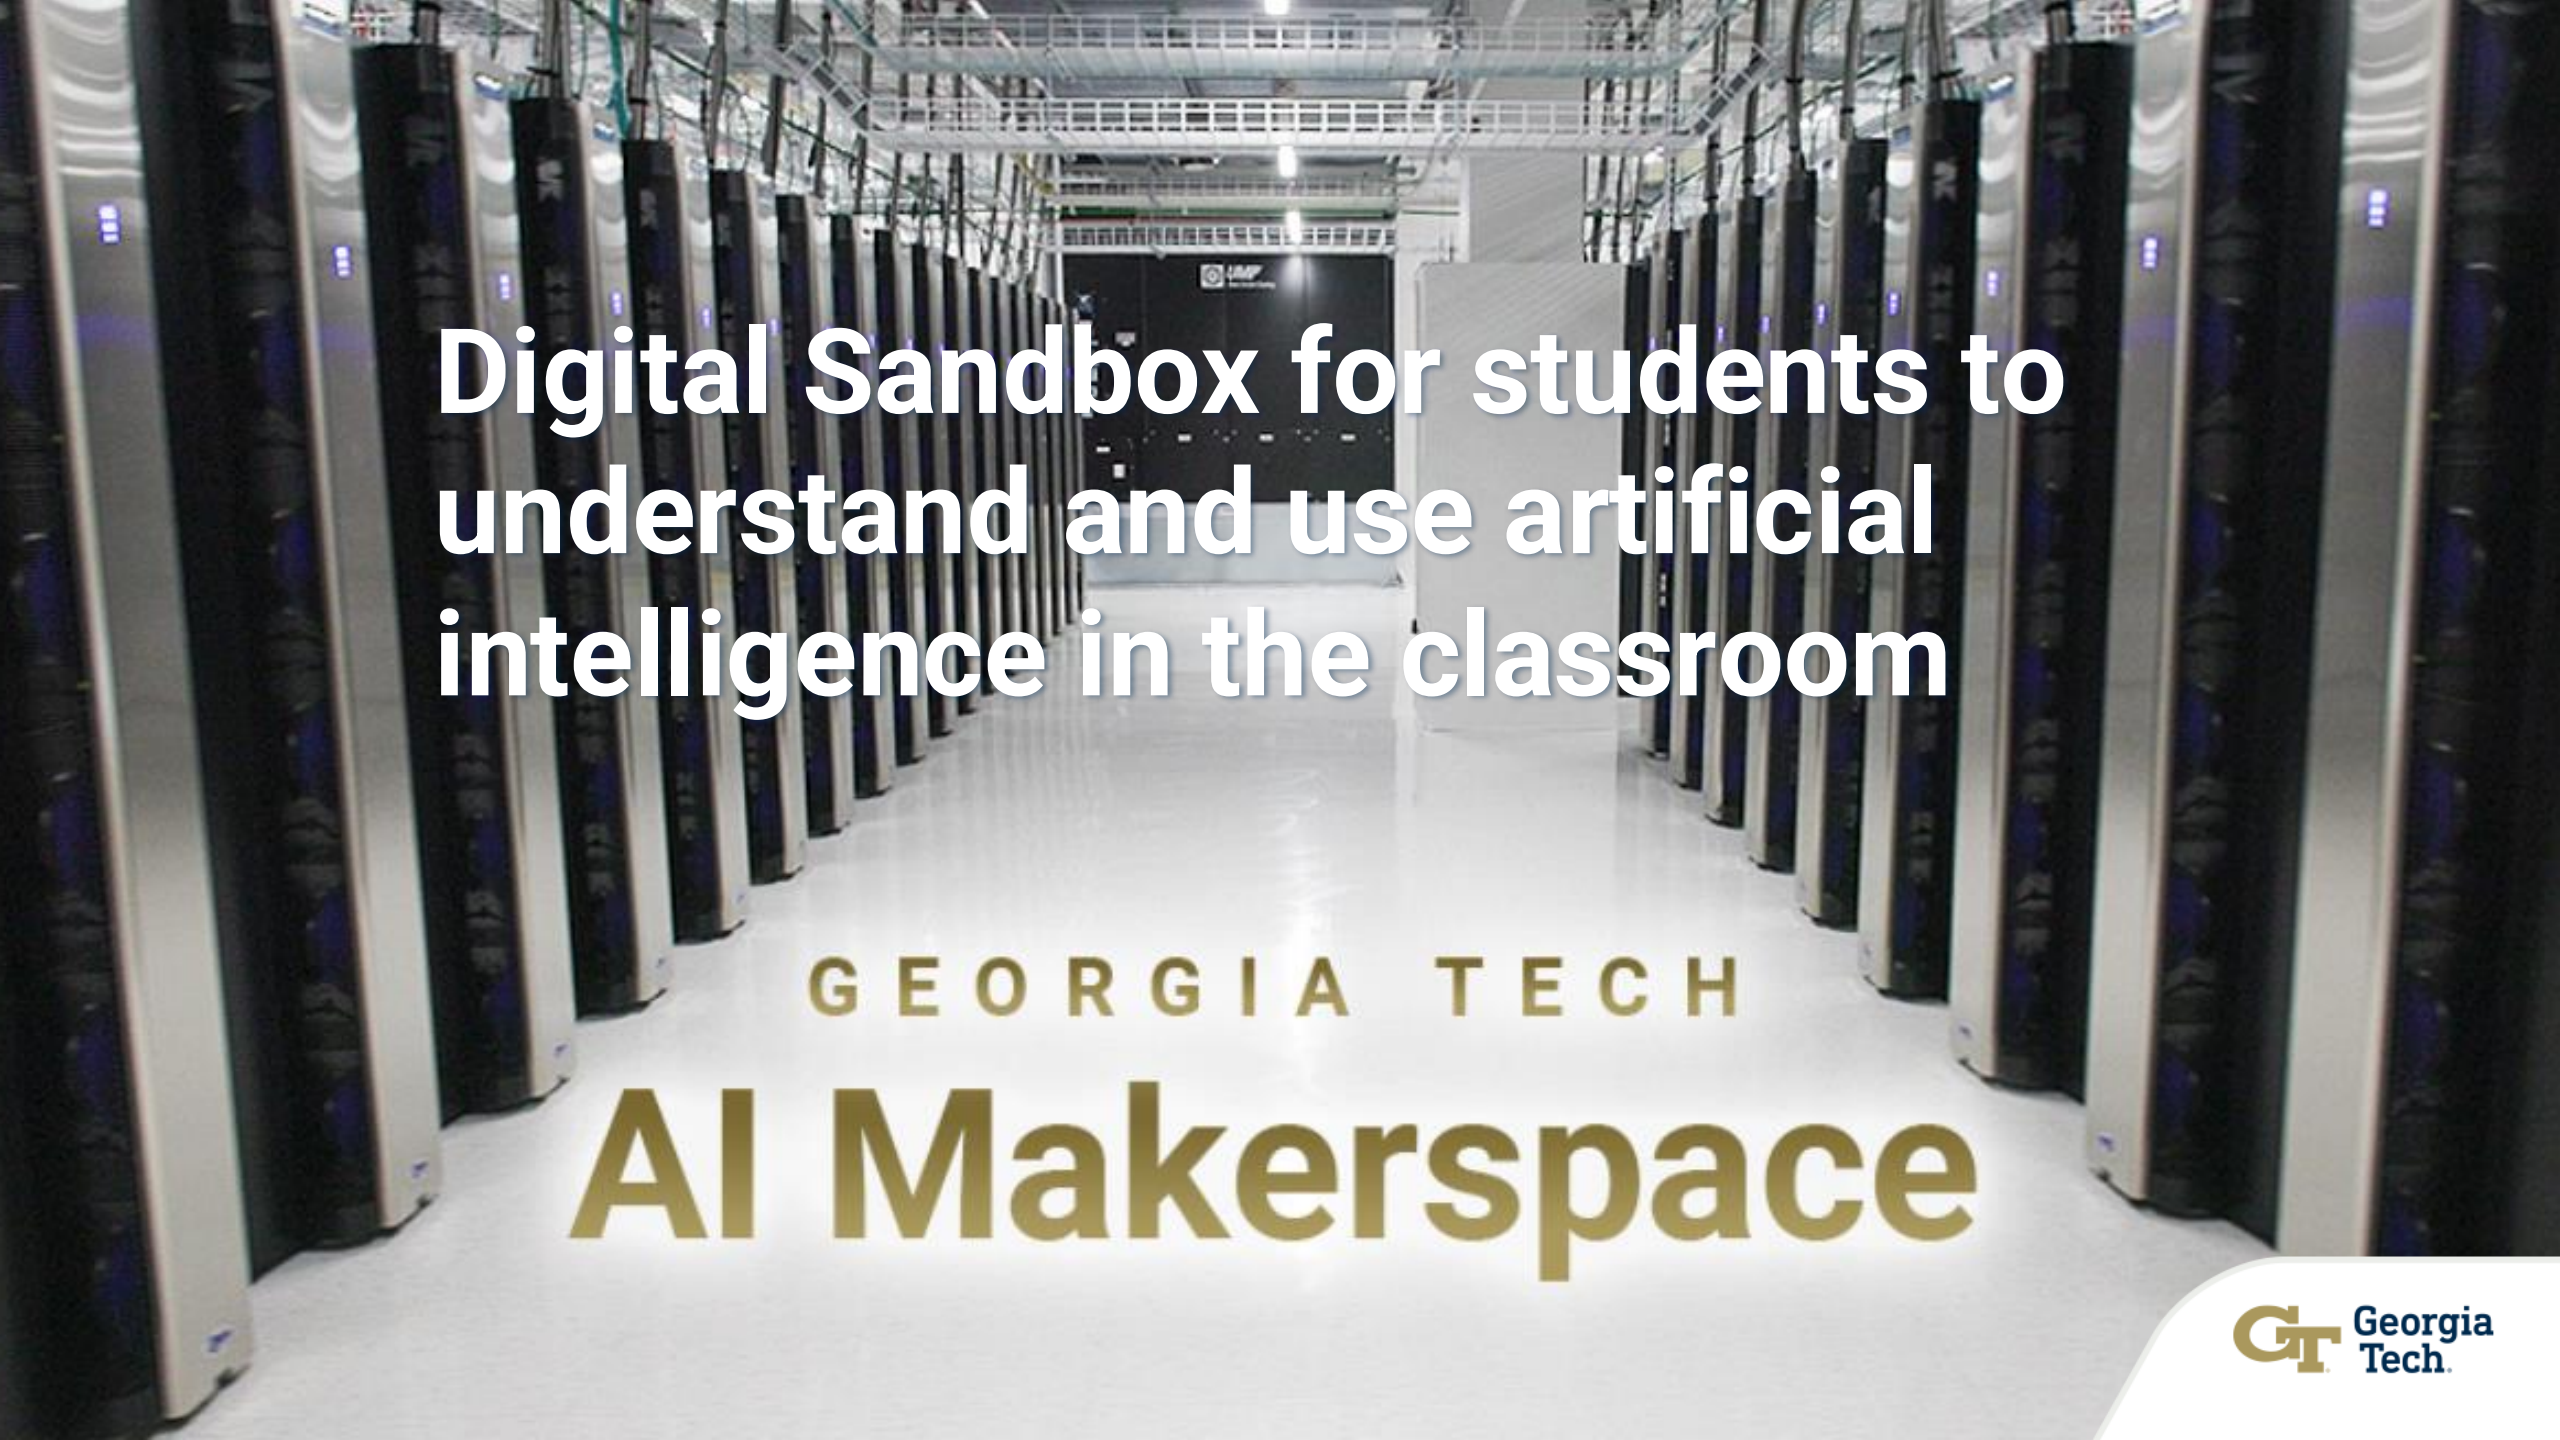
Digital Sandbox for students to understand and use artificial intelligence in the classroom
GEORGIA TECH
AI Makerspace

Georgia Tech AI Makerspace

- Officially released on April 10th , 2024
- The first AI computing resources designed for student use within the nation
- Collaborated with Nvidia, GT College of Engineering and PACE computing center
- Phase I contains 20 8-way HGX H100 boxes, makes total 120GPUs
- Four courses are using the system at Fall 2024
- Pending Phase II contains 18 8-way HGX H200 boxes, adds 144 more powerful GPUs, system is currently under testing

Georgia Tech

Timeline for AI Makerspace Delivery

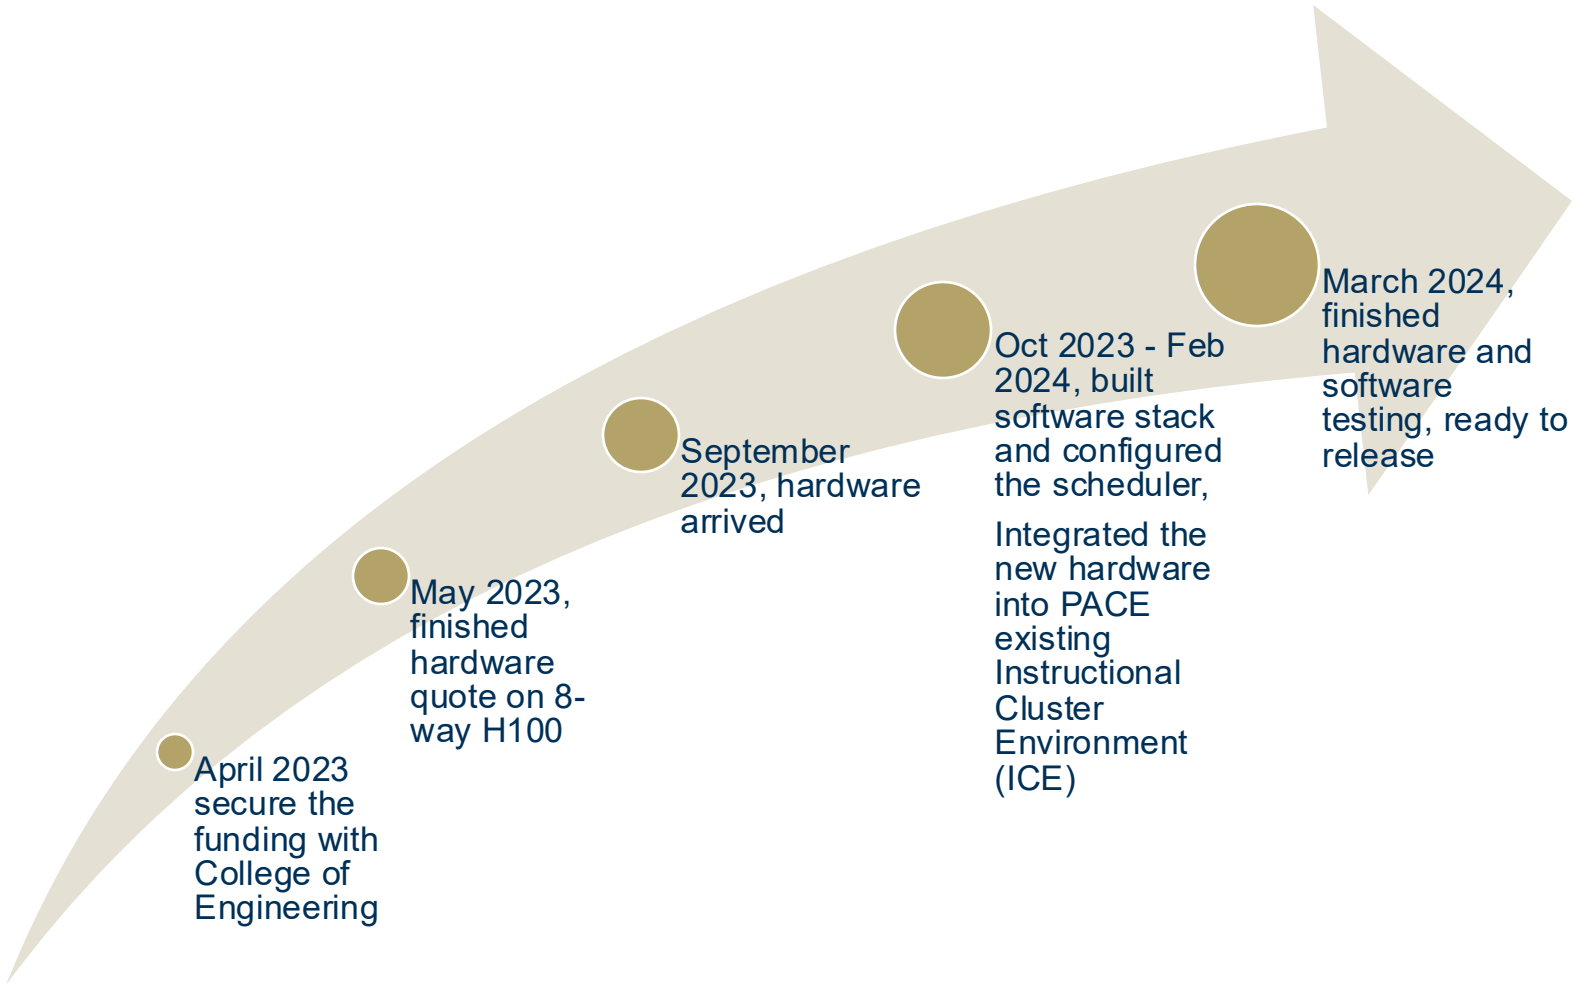

Georgia Tech

PACE - Partnership for an Advanced Computing Environments

- Georgia Tech's PACE center (http://pace.gatech.edu) provides scalable HPC and instructional resources for Georgia Tech researchers and students
- Supports multiple clusters:
  - Phoenix – research cluster
  - Hive – NSF MRI resource
  - Firebird – CUI/ITAR-complaint research cluster
  - ICE – Instructional cluster for courses and educational workshops
- Multi-Team structure:
  - Cyber Infrastructure
  - Architecture & Platforms
  - Research Computing Facilitation & Customer Engagement
  - Research Software Engineering

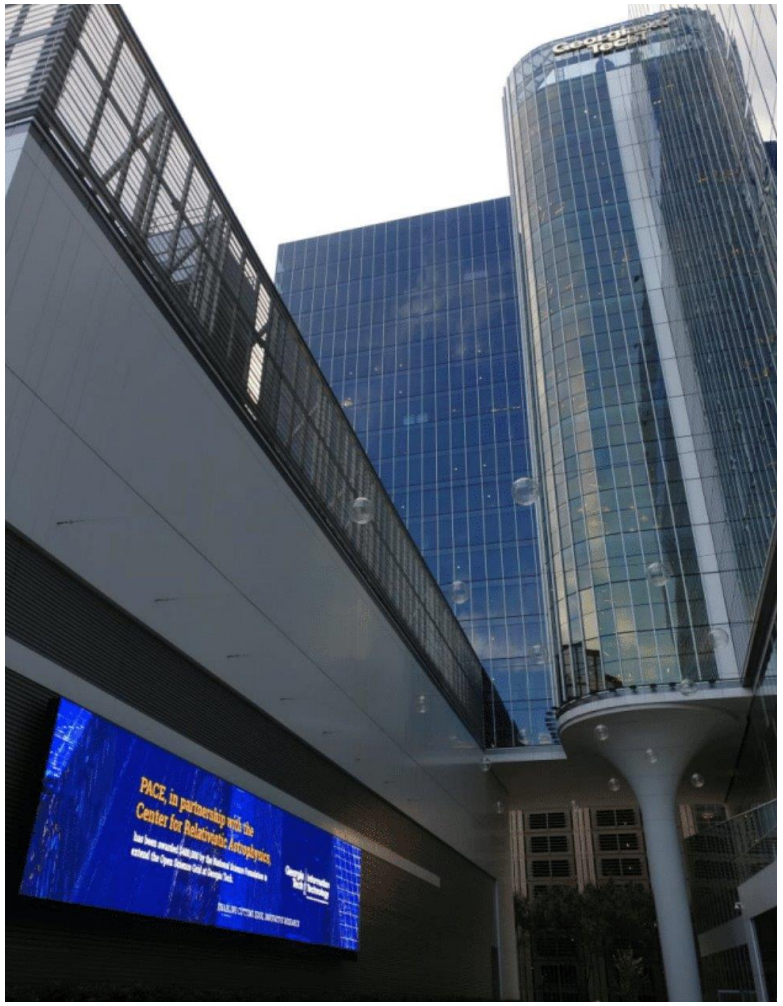

Georgia Tech

PACE Mission

Serve & Empower

Research

General Research Compute (Phoenix) : Any GT Faculty

Controlled, Unclassified Information (CUI) Research Compute (Firebird): As needed

NSF-MRI cluster (HIVE): Limited faculty, 20% allocated to **ACCESS**

Teaching & Learning

Instructional Cluster Environments (CoC-ICE & PACE-ICE): Dedicated to Scientific Computing Instruction

Technical Seminars & Tutorials

Democratization

EVPR-PACESHIP: $200k student scholarships

Open Science Grid (Buzzard): Funded by an NSF CC* award

ACCESS

Free Tier (Compute & Storage)

Outreach & Engagement

Georgia Tech

An ecosystem to support students in the different level of AI Knowledge and Background

Georgia Tech

Challenges

Students/Instructors are lacking experience on using HPC systems

Lack of an interactive environment to develop new course material

Large training dataset I/O performance

CPU and GPU workload efficiency

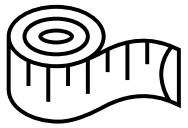

Portability of Nvidia NGC containers to HPC systems

Georgia Tech

Lowering the Access Barrier through Open OnDemand

- Funded by NSF and developed by Ohio Supercomputing Center
- Provides a web-based job submission and Jupyter Notebook interface
- Gives the freedom to add new and customized application and hides the details of resource requests
- ECE 2806 – Foundations of AI course is released on PACE Open OnDemand

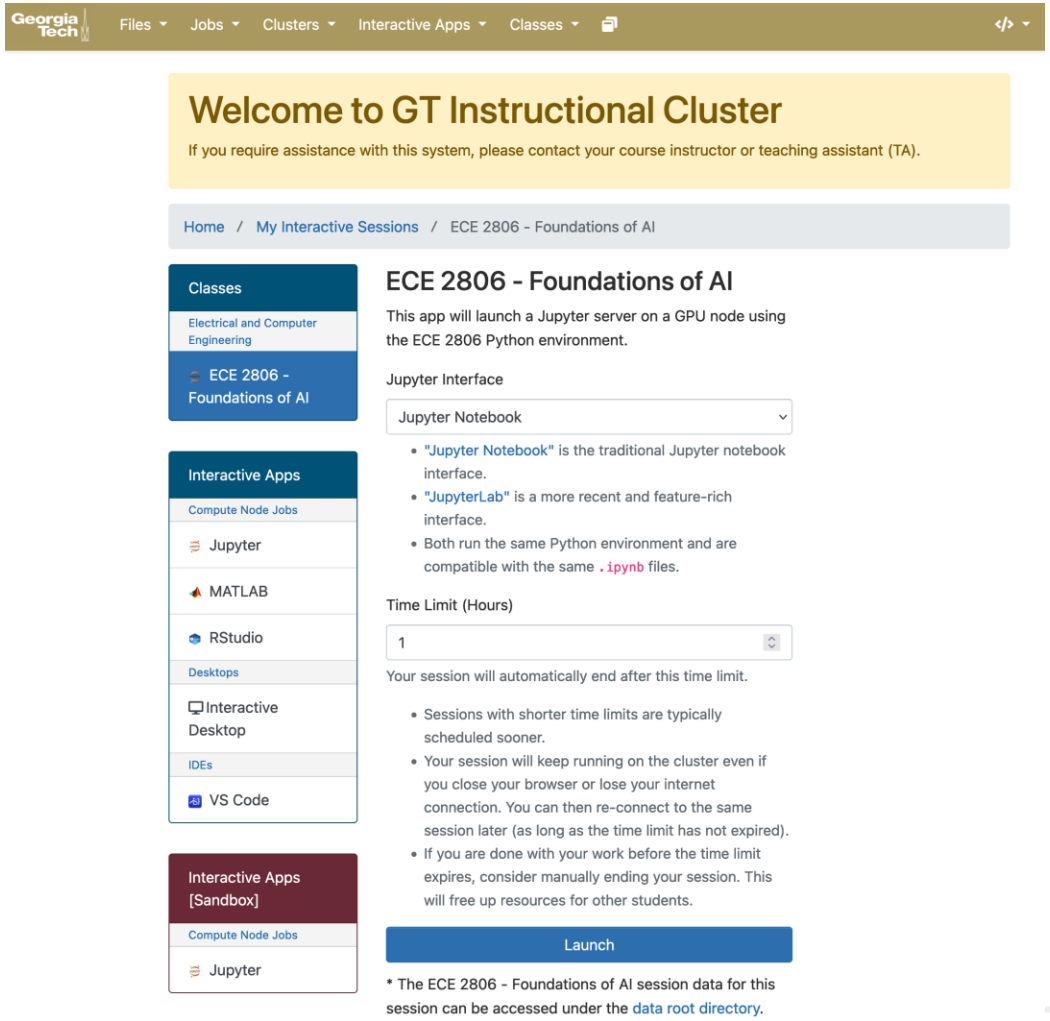

Georgia Tech

Container Support through Apptainer

- Docker is not supported in the PACE cluster and in many research computing centers
  - Podman does provide Docker-like support but requires using newer OS support features and has some NFS limitations
- Apptainer is the primary alternative to Docker on HPC systems
  - allows unprivileged users to use containers
  - prohibits escalation within the container
- To support AI/ML workflows:
  - Nvidia's NGC docker containers are converted into Apptainer containers for common scenarios, e.g. PyTorch, Tensorflow
  - Customized containers are built from base images from Nvidia
  - Integrate the container to Open OnDemand Jupyter Notebook interface
- We are in the process to enable container self-service in which everyone can build and run containers on the AI Makerspace

Large AI Training Dataset Support

- AI training datasets usually contain many small files and a large total size
- PACE provides a central location to store AI datasets to avoid duplication issues for these common datasets
- In order to determine the best storage location, we compared the I/O performance across all available filesystems we host
  - e.g. Lustre, pNFS over RDMA, NFD over TCP, Local disk
- The comparison was done across different data sizes, data formats and filesystems

[1] Exploring Research Dataset-Sharing Strategies for Concurrent AI Workflows
https://dl.acm.org/doi/10.1145/3626203.3670597 (Best student paper at PEARC'24)

Conclusions

PACE Accomplishments with the AI Makerspace Include:

- Enablement for instructors to create trainings with more complex real-world problems
- Support for student teams' senior design projects in multiple semesters
- Hosting for training events (e.g. Nvidia Hackathon) for internal and external users including high school students
- Enhancement of vendor relationships – Nvidia and Penguin

Lessons Learned

- PACE experienced a shortage of RSE-related support resources, leading to a challenging deployment
- We should also engage stakeholders ahead of time to ensure the usage of resources is fully understood

The AI Makerspace offers students cutting-edge GPU capabilities in a classroom setting!

Georgia Tech

AI Makerspace Phase II

- Hardware (with 18 8-way HGX H200 nodes) is under testing which will at least double the computing capacities
- Work to establish the student governance to make Makerspace a fully student-ran resource
- Enhance user services to accommodate the broad access to the Makerspace
- Integrate more vendor provided software solutions

Questions

Dr. Fang (Cherry) Liu fang.liu@gatech.edu

Georgia Tech